

# Class Imbalance

My business goal is identifying whether the company will bankrupt or not in next 1-5 years.

I will use Ensemble Method 'XGBoost', eXtreme Gradient Boosting, for classification.

I will focus on the performance of 'recall' metric in order to minimize false negatives. Besides, I will also keep an eye on 'precision', 'f1', 'accuracy, and 'AUC' metrics.

**In this botebook, I will work on the class imbalance issue on the dataset.**

## Load Libraries

```
In [16]: # Import base libraries
import pandas as pd
import numpy as np
from scipy.io import arff

import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
from xgboost import XGBClassifier

from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import roc_auc_score, roc_curve, auc
from sklearn.metrics import precision_score, recall_score, accuracy_score,
from sklearn.utils import class_weight

from functions import *

from datetime import datetime

import warnings
warnings.filterwarnings('ignore')
```

## Load Data

There are five data files:

- \* data1, 1year.arff
- \* data2, 2year.arff
- \* data3, 3year.arff
- \* data4, 4year.arff
- \* data5, 5year.arff

Initially, I will use 'data3' to find the best approach to handle the class imbalance.

No cleaning applied to data. XGBoost Classifier can handle the missing values.

```
In [23]: # Load data
data3 = arff.loadarff('data/3year.arff')
df3 = pd.DataFrame(data3[0])

# Change label/class type to binary
df3['class'] = df3['class'].astype('int64')

df3.shape
```

```
Out[23]: (10503, 65)
```

## Pre-process

```
In [24]: # Assign target and predictor
y = df3['class']
X = df3.drop('class', axis=1)

# Separate data into train and test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, r

# Scale/Normalize the predictor variables
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
y_train = y_train.to_numpy()
y_test = y_test.to_numpy()

print('X_train shape = ', X_train.shape)
print('y_train shape = ', y_train.shape)
print('X_test shape = ', X_test.shape)
print('y_test shape = ', y_test.shape)
```

```
X_train shape = (8402, 64)
y_train shape = (8402,)
X_test shape = (2101, 64)
y_test shape = (2101,)
```

## Baseline Model

```
In [19]: # Baseline Model

xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
}

baseline_model = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbP
```

Data 3 Classification Report:

Training Data:

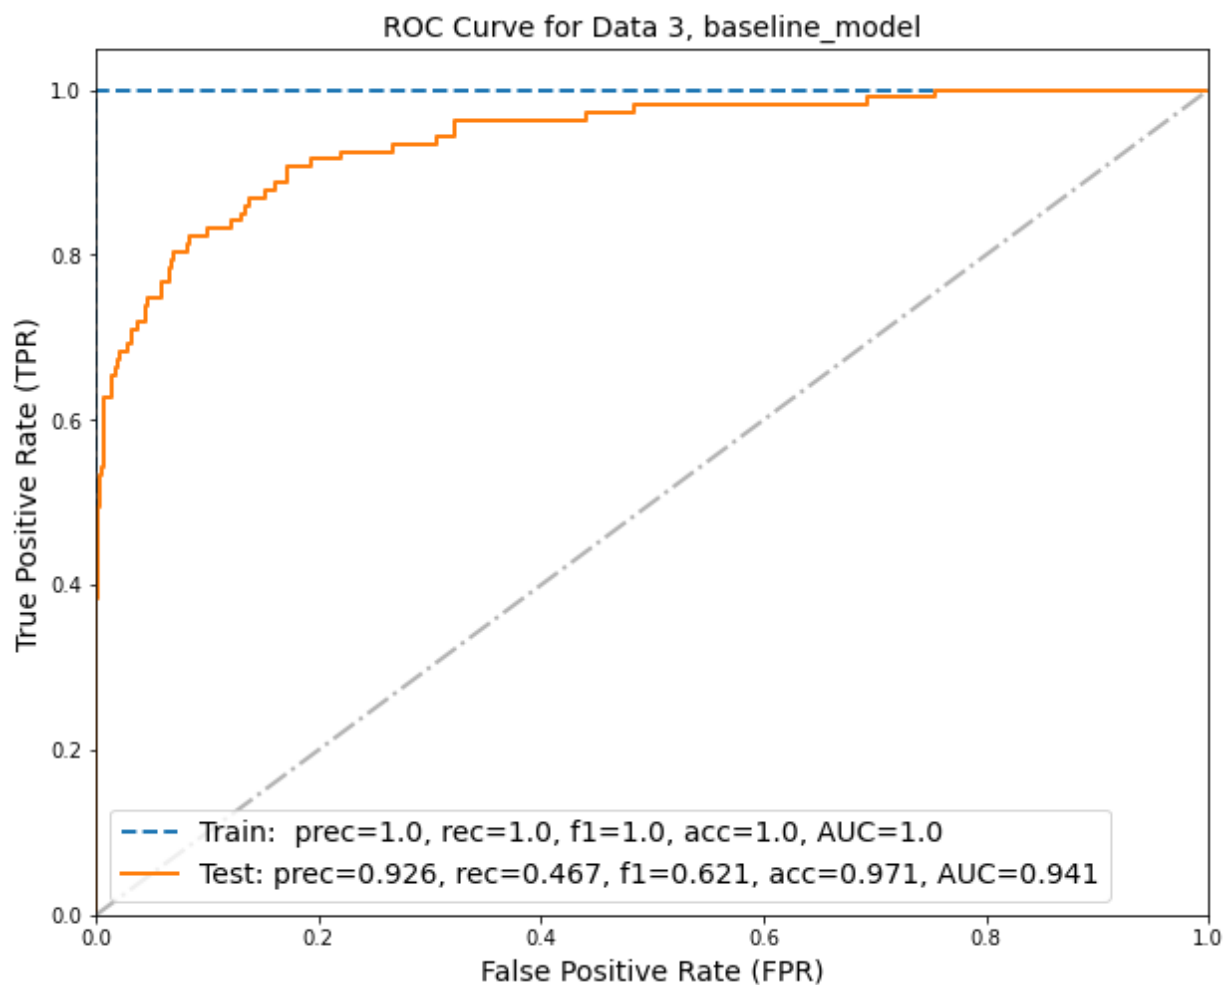
	precision	recall	f1-score	support
0	1.00	1.00	1.00	8014
1	1.00	1.00	1.00	388
accuracy			1.00	8402
macro avg	1.00	1.00	1.00	8402
weighted avg	1.00	1.00	1.00	8402

Testing Data:

	precision	recall	f1-score	support
0	0.97	1.00	0.98	1994
1	0.93	0.47	0.62	107
accuracy			0.97	2101
macro avg	0.95	0.73	0.80	2101
weighted avg	0.97	0.97	0.97	2101

```
In [6]: #Plot ROC curve
```

```
plot_ROC(3, X_train, y_train, X_test, y_test, baseline_model, 'baseline_mod
```



```
In [7]: model_1.get_params
```

```
Out[7]: <bound method XGBModel.get_params of XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bynode=1, colsample_bytree=1, eval_metric='logloss',
               gamma=0, gpu_id=-1, importance_type='gain',
               interaction_constraints='', learning_rate=0.300000012,
               max_delta_step=0, max_depth=6, min_child_weight=1, missing=nan,
               monotone_constraints='()', n_estimators=100, n_jobs=4,
               num_parallel_tree=1, random_state=42, reg_alpha=0, reg_lambda=1,
               scale_pos_weight=1, subsample=1, tree_method='exact',
               validate_parameters=1, verbosity=None)>
```

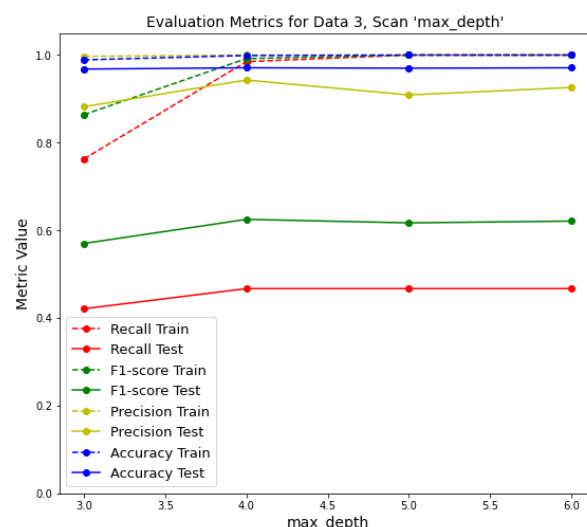
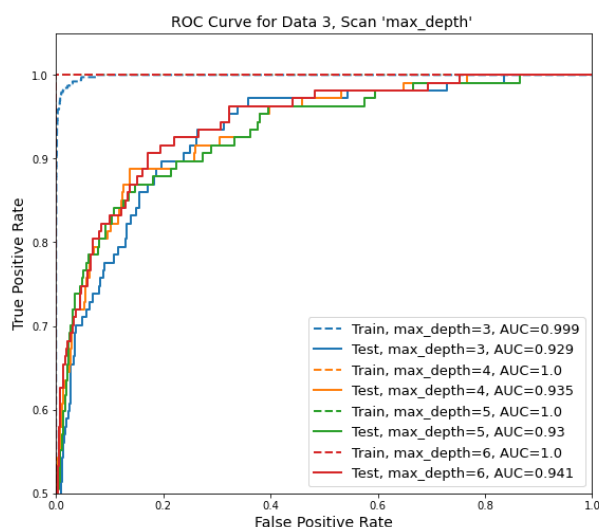
```
In [33]: #Scan max_depth
```

```
xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
}

scanParam = 'max_depth'

scanList = [3, 4, 5, 6]

result_sw_depth = scan_xgb_ROC_metrics(3, X_train, y_train, X_test, y_test,
```



## Comments

- According to testing data results, the model performance in identifying the non-bankruptcy companies (class 0) is very good.
- However, the identification of the bankruptcy companies (class 1) is not that good (low recall and f1)
- The class imbalance is probably the reason for the class 0/1 performance difference.
- There is large overfitting.

# Class Imbalance

There are two approaches to deal with the class imbalance:

- 'sample\_weight' parameter when training the data
- 'scale\_pos\_weight' parameter when initiating the classifier

## sample\_weight

In [24]: *# Whole data 3*

```
print('Whole Data3')
d3_class_weights = class_weight.compute_class_weight(class_weight='balanced', classes=d3_classes, samples=d3_samples)
print(d3_class_weights)

d3_class_weights_ratio = d3_class_weights[1]/d3_class_weights[0]
print('weights_ratio', d3_class_weights_ratio)
print('Squared Root of weights_ratio', np.sqrt(d3_class_weights_ratio))
```

```
Whole Data3
[ 0.52473022 10.60909091]
weights_ratio 20.21818181818182
Squared Root of weights_ratio 4.496463256625347
```

In [28]: *#Training sample*

```
print('Training sample, Data3')
d3_class_weights_train = class_weight.compute_class_weight(class_weight='balanced', classes=d3_classes, samples=d3_train_samples)
print(d3_class_weights_train)
d3_class_weights_train_ratio = d3_class_weights_train[1]/d3_class_weights_train[0]
print('weights_ratio', d3_class_weights_train_ratio)
print('Squared Root of weights_ratio', np.sqrt(d3_class_weights_train_ratio))

d3_class_weights_train_array = class_weight.compute_sample_weight(class_weight=d3_class_weights_train, samples=d3_train_samples)
print('Class weights per entry:', d3_class_weights_train_array)

# Training weights are used for data training.
```

```
Training sample, Data3
weights_ratio 20.65463917525773
Squared Root of weights_ratio 4.544737525452678
Class weights per entry: [0.52420764 0.52420764 0.52420764 ... 0.52420764
0.52420764 0.52420764]
```

```
In [31]: # sample_weight applied

xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
}

xgb_model_sw = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbPar
```

Sample weights are used!

-----

### Data 3 Classification Report:

#### Training Data:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8014
1	1.00	1.00	1.00	388
accuracy			1.00	8402
macro avg	1.00	1.00	1.00	8402
weighted avg	1.00	1.00	1.00	8402

#### Testing Data:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	1994
1	0.87	0.54	0.67	107
accuracy			0.97	2101
macro avg	0.92	0.77	0.83	2101
weighted avg	0.97	0.97	0.97	2101

```
In [35]: #Scan max_depth
#sample_weight applied

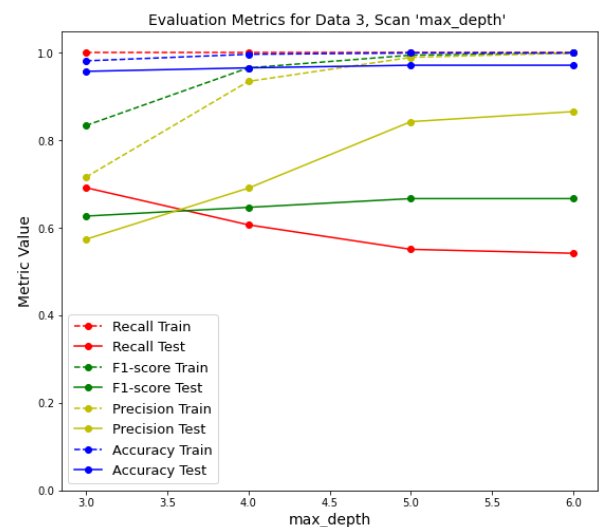

xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
}

scanParam = 'max_depth'

scanList = [3, 4, 5, 6]

result_sw_depth = scan_xgb_ROC_metrics(3, X_train, y_train, X_test, y_test,

Sample weights are used!
-----
```



## Comments

- Model performance improved, Recall increased.
- Overfitting continues.
- sample\_weight is more effective at lower max\_depth (3 or 4)

## scale\_pos\_weight

My goal is find the optimum 'scale\_pos\_weight' value which creates good recall and f1 value.

The model performance is largely effected by the max\_depth. So I will scan the scale\_pos\_weight at several max\_depth.



```
In [29]: #Check Class Imbalance
df3['class'].value_counts()
```

```
Out[29]: 0    10008
         1     495
         Name: class, dtype: int64
```

```
In [30]: #Check Class Imbalance, Normalized
df3['class'].value_counts(normalize=True)
```

```
Out[30]: 0    0.952871
         1    0.047129
         Name: class, dtype: float64
```

```
In [15]: # Explore class counts
val_counts = df3['class'].value_counts()
ratio_imbalance = val_counts[0]/val_counts[1]
sqrt_ratio_imbalance = np.sqrt(val_counts[0]/val_counts[1])

print('imbalance ratio:', ratio_imbalance)
print('sqrt of imbalance ratio:', sqrt_ratio_imbalance)

# The values are very similar for train/test/whole datasets.
```

```
imbalance ratio: 20.21818181818182
sqrt of imbalance ratio: 4.496463256625347
```

```
In [22]: # Explore class counts, Training sample
val_counts = pd.DataFrame(y_train, columns=['class'])['class'].value_counts
ratio_imbalance = val_counts[0]/val_counts[1]
sqrt_ratio_imbalance = np.sqrt(val_counts[0]/val_counts[1])

print('imbalance ratio:', ratio_imbalance)
print('sqrt of imbalance ratio:', sqrt_ratio_imbalance)
```

```
imbalance ratio: 20.65463917525773
sqrt of imbalance ratio: 4.544737525452678
```

```

In [37]: #Scan scale_pos_weight,
# max_depth=6 (default)


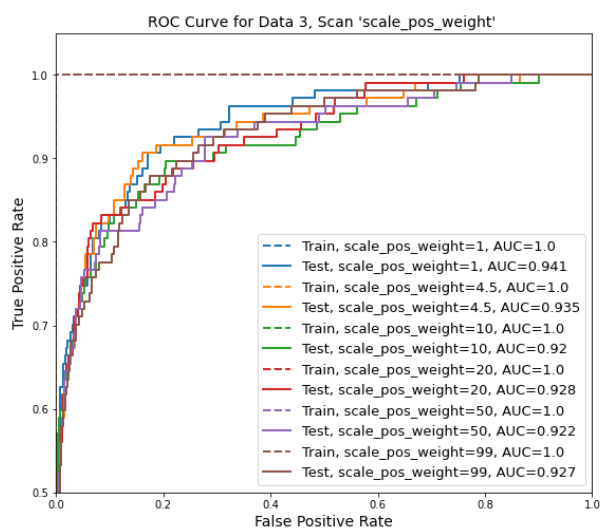
xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
}

scanParam = 'scale_pos_weight'

scanList = [1, 4.5, 10, 20, 50, 99]
#scanList = [1, 4.5]

result_spw = scan_xgb_ROC_metrics(3, X_train, y_train, X_test, y_test, xgbP

```



```
In [51]: # max_depth=6, scale_pos_weight=10
```

```
xgbParams = {  
    'eval_metric': 'logloss',  
    'random_state': 42,  
    'max_depth': 6,  
    'scale_pos_weight': 10,  
}
```

```
model_spw10_depth6 = xgb_model_report(3, X_train, y_train, X_test, y_test,
```

Data 3 Classification Report:

Training Data:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8014
1	1.00	1.00	1.00	388
accuracy			1.00	8402
macro avg	1.00	1.00	1.00	8402
weighted avg	1.00	1.00	1.00	8402

Testing Data:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	1994
1	0.88	0.56	0.69	107
accuracy			0.97	2101
macro avg	0.93	0.78	0.84	2101
weighted avg	0.97	0.97	0.97	2101

```

In [40]: #Scan scale_pos_weight,
# max_depth=5


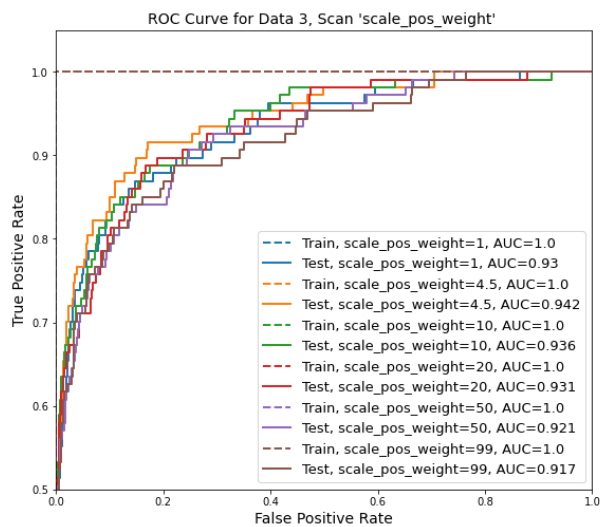
xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 5
}

scanParam = 'scale_pos_weight'

scanList = [1, 4.5, 10, 20, 50, 99]
#scanList = [1, 4.5]

result_spw_depth5 = scan_xgb_ROC_metrics(3, X_train, y_train, X_test, y_test)

```



```
In [52]: # max_depth=5, scale_pos_weight=50
```

```
xgbParams = {  
    'eval_metric': 'logloss',  
    'random_state': 42,  
    'max_depth': 5,  
    'scale_pos_weight': 50,  
}
```

```
model_spw50_depth5 = xgb_model_report(3, X_train, y_train, X_test, y_test,
```

Data 3 Classification Report:

Training Data:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8014
1	0.99	1.00	0.99	388
accuracy			1.00	8402
macro avg	0.99	1.00	1.00	8402
weighted avg	1.00	1.00	1.00	8402

Testing Data:

	precision	recall	f1-score	support
0	0.98	0.99	0.98	1994
1	0.76	0.56	0.65	107
accuracy			0.97	2101
macro avg	0.87	0.78	0.81	2101
weighted avg	0.97	0.97	0.97	2101

```

In [41]: #Scan scale_pos_weight,
# max_depth=4

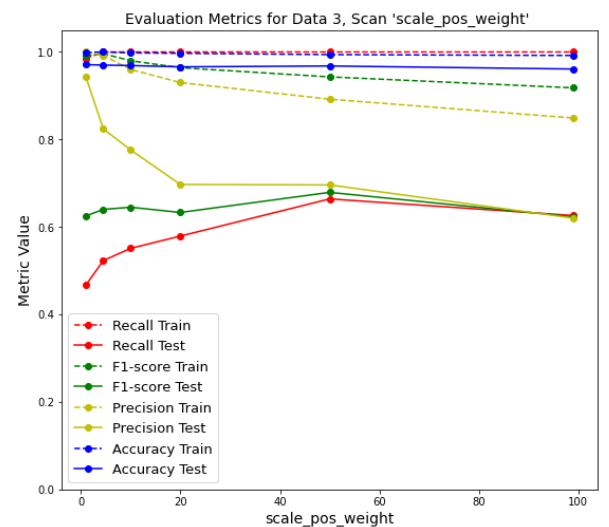
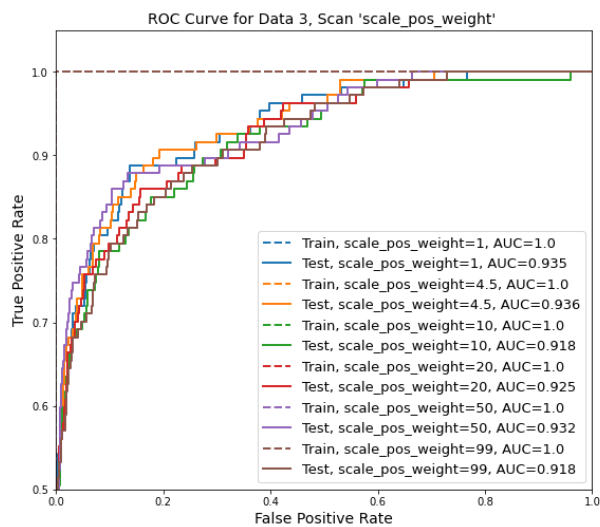
xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 4
}

scanParam = 'scale_pos_weight'

scanList = [1, 4.5, 10, 20, 50, 99]
#scanList = [1, 4.5]

result_spw_depth4= scan_xgb_ROC_metrics(3, X_train, y_train, X_test, y_test

```



```
In [53]: # max_depth=4, scale_pos_weight=50
```

```
xgbParams = {  
    'eval_metric': 'logloss',  
    'random_state': 42,  
    'max_depth': 4,  
    'scale_pos_weight': 50,  
}
```

```
model_spw50_depth4 = xgb_model_report(3, X_train, y_train, X_test, y_test,
```

Data 3 Classification Report:

Training Data:

	precision	recall	f1-score	support
0	1.00	0.99	1.00	8014
1	0.89	1.00	0.94	388
accuracy			0.99	8402
macro avg	0.95	1.00	0.97	8402
weighted avg	1.00	0.99	0.99	8402

Testing Data:

	precision	recall	f1-score	support
0	0.98	0.98	0.98	1994
1	0.70	0.66	0.68	107
accuracy			0.97	2101
macro avg	0.84	0.82	0.83	2101
weighted avg	0.97	0.97	0.97	2101

```

In [42]: #Scan scale_pos_weight,
# max_depth=3

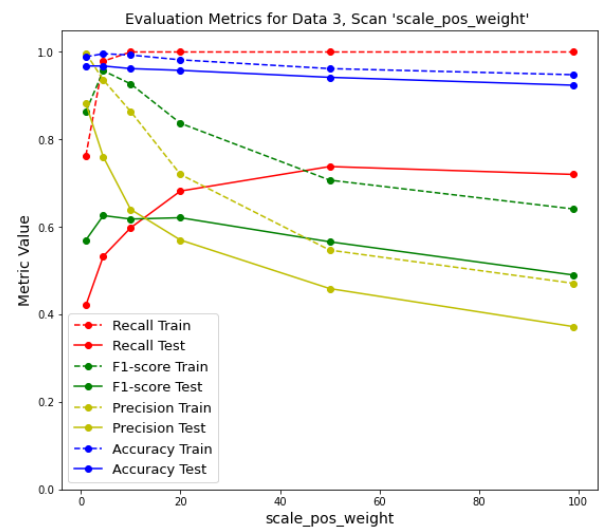

xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 3
}

scanParam = 'scale_pos_weight'

scanList = [1, 4.5, 10, 20, 50, 99]
#scanList = [1, 4.5]

result_spw_depth3= scan_xgb_ROC_metrics(3, X_train, y_train, X_test, y_test

```





```
In [55]: # max_depth=3, scale_pos_weight=20
```

```
xgbParams = {  
    'eval_metric': 'logloss',  
    'random_state': 42,  
    'max_depth': 3,  
    'scale_pos_weight': 20,  
}
```

```
model_spw20_depth3 = xgb_model_report(3, X_train, y_train, X_test, y_test,
```

Data 3 Classification Report:

Training Data:

	precision	recall	f1-score	support
0	1.00	0.98	0.99	8014
1	0.72	1.00	0.84	388
accuracy			0.98	8402
macro avg	0.86	0.99	0.91	8402
weighted avg	0.99	0.98	0.98	8402

Testing Data:

	precision	recall	f1-score	support
0	0.98	0.97	0.98	1994
1	0.57	0.68	0.62	107
accuracy			0.96	2101
macro avg	0.78	0.83	0.80	2101
weighted avg	0.96	0.96	0.96	2101

```

In [56]: #Compare models with optimum scale_pos_weight

model_list = [model_spw10_depth6, model_spw50_depth5, model_spw50_depth4, m

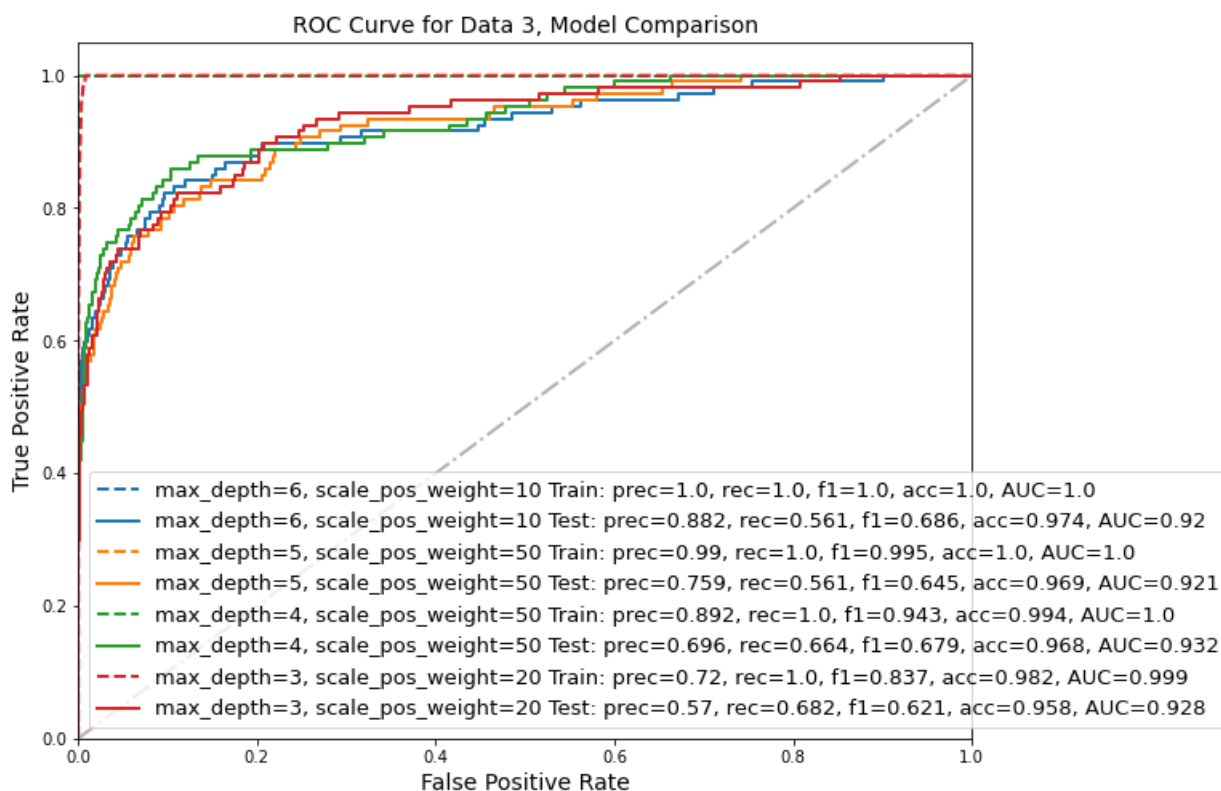
model_names_list = [ 'max_depth=6, scale_pos_weight=10',
                      'max_depth=5, scale_pos_weight=50',
                      'max_depth=4, scale_pos_weight=50',
                      'max_depth=3, scale_pos_weight=20',
                      ]

compare_models(3, X_train, y_train, X_test, y_test, model_list, model_names

```

Out[56]:

	precision	recall	f1	accuracy	auc
Params					
max_depth=6, scale_pos_weight=10 Train	1.000	1.000	1.000	1.000	1.000
Test	0.882	0.561	0.686	0.974	0.920
max_depth=5, scale_pos_weight=50 Train	0.990	1.000	0.995	1.000	1.000
Test	0.759	0.561	0.645	0.969	0.921
max_depth=4, scale_pos_weight=50 Train	0.892	1.000	0.943	0.994	1.000
Test	0.696	0.664	0.679	0.968	0.932
max_depth=3, scale_pos_weight=20 Train	0.720	1.000	0.837	0.982	0.999
Test	0.570	0.682	0.621	0.958	0.928



### **Comments**

- The model performance improves compared to the baseline model (Model 1) where default `scale_pos_weight=1`.
- The optimum value may be different when other parameters change, for instance `max_depth`.
- The optimum `scale_pos_weight`:
  - `scale_pos_weight=10` at `max_depth=6`
  - `scale_pos_weight=50` at `max_depth=5`
  - `scale_pos_weight=50` at `max_depth=4`
  - `scale_pos_weight=20` at `max_depth=3`
- Overfitting observed.
- The optimum value may be different when another data set is used. It may require additional parameter tuning.

### **`class_weight + scale_pos_weight`**

```

In [57]: #Scan scale_pos_weight
# sample_weight applied
#max_depth=6

xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
}

scanParam = 'scale_pos_weight'


scanList = [1, 4.5, 10, 20, 50, 99]
#scanList = [1, 4.5]

result_spw_sw_depth6 = scan_xgb_ROC_metrics(3, X_train, y_train, X_test, y_

```

Sample weights are used!

-----



```

In [58]: #Scan scale_pos_weight
# sample_weight applied
# max_depth=5

xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 5,
}

scanParam = 'scale_pos_weight'



scanList = [1, 4.5, 10, 20, 50, 99]
#scanList = [1, 4.5]

result_spw_sw_depth5 = scan_xgb_ROC_metrics(3, X_train, y_train, X_test, y_

```

Sample weights are used!

-----



```
In [60]: #Scan scale_pos_weight
# sample_weight applied
# max_depth=4



xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 4,
}

scanParam = 'scale_pos_weight'

scanList = [1, 4.5, 10, 20, 50, 99]
#scanList = [1, 4.5]

result spw sw depth4 = scan xgb ROC metrics(3, X_train, y_train, X_test, y_test)
```

Sample weights are used!



```

In [59]: #Scan scale_pos_weight
# sample_weight applied
# max_depth=3

xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 3,
}

scanParam = 'scale_pos_weight'

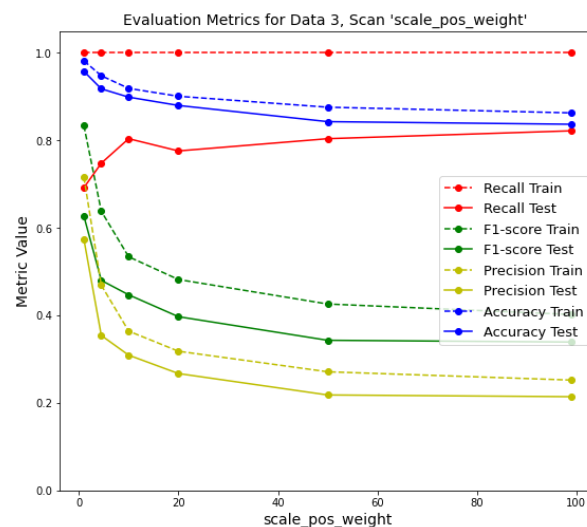

scanList = [1, 4.5, 10, 20, 50, 99]
#scanList = [1, 4.5]

result_spw_sw_depth3 = scan_xgb_ROC_metrics(3, X_train, y_train, X_test, y_

```

Sample weights are used!

-----



In [14]: *#Compare balance methods at max\_depth=6*

```
xgbParams1 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 6,
}

mod1 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams1, 'X

xgbParams2 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 6,
    'scale_pos_weight': 10,
}

mod2 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams2, 'X

xgbParams3 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 6,
    'scale_pos_weight': 20,
}

mod3 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams3, 'X

xgbParams4 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 6,
    'scale_pos_weight': 50,
}

mod4 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams4, 'X

model_list = [mod1, mod2, mod3, mod4]

model_names_list = ['sample_weight', 'scale_pos_weight=10', 'sample_weight

compare_models(3, X_train, y_train, X_test, y_test, model_list, model_names
```

Sample weights are used!

-----

Sample weights are used!

-----

Sample weights are used!


-----

Out[14]:

precision	recall	f1	accuracy	auc
-----------	--------	----	----------	-----



Params	precision	recall	f1	accuracy	auc
Params					
sample_weight Train	1.000	1.000	1.000	1.000	1.000
Test	0.866	0.542	0.667	0.972	0.931
scale_pos_weight=10 Train	1.000	1.000	1.000	1.000	1.000
Test	0.882	0.561	0.686	0.974	0.920
sample_weight + scale_pos_weight=20 Train	0.990	1.000	0.995	1.000	1.000
Test	0.677	0.607	0.640	0.965	0.916
sample_weight + scale_pos_weight=50 Train	0.985	1.000	0.992	0.999	1.000
Test	0.663	0.607	0.634	0.964	0.914





```

In [8]: #Compare balance methods at max_depth=5

xgbParams1 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 5,
}

mod1 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams1, 'X

xgbParams2 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 5,
    'scale_pos_weight': 50,
}

mod2 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams2, 'X

xgbParams3 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 5,
    'scale_pos_weight': 20,
}

mod3 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams3, 'X

model_list = [mod1, mod2, mod3]

model_names_list = ['sample_weight', 'scale_pos_weight=50', 'sample_weight

compare_models(3, X_train, y_train, X_test, y_test, model_list, model_names

```

Sample weights are used!

-----

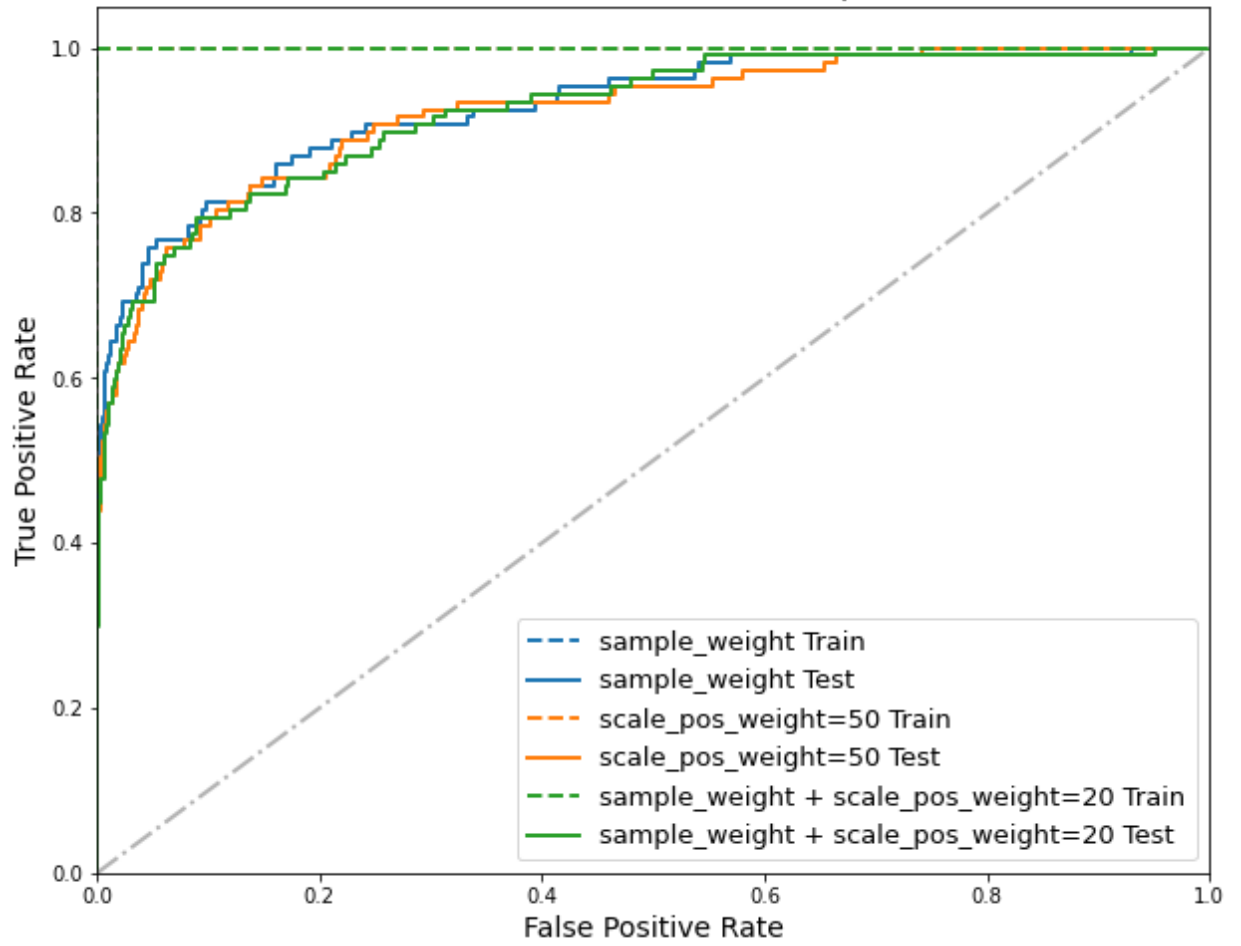
Sample weights are used!

-----

Out[8]:

		precision	recall	f1	accuracy	auc
Params						
sample_weight Train		0.990	1.000	0.995	1.000	1.000
	Test	0.843	0.551	0.667	0.972	0.927
scale_pos_weight=50 Train		0.990	1.000	0.995	1.000	1.000
	Test	0.759	0.561	0.645	0.969	0.921
sample_weight + scale_pos_weight=20 Train		0.946	1.000	0.972	0.997	1.000
	Test	0.624	0.636	0.630	0.962	0.921

ROC Curve for Data 3, Model Comparison



In [13]: *#Compare balance methods at max\_depth=4*

```
xgbParams1 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 4,
}

mod1 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams1, 'X

xgbParams2 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 4,
    'scale_pos_weight': 50,
}

mod2 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams2, 'X

xgbParams3 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 4,
    'scale_pos_weight': 4.5,
}

mod3 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams3, 'X

xgbParams4 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 4,
    'scale_pos_weight': 50,
}

mod4 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams4, 'X

model_list = [mod1, mod2, mod3, mod4]

model_names_list = ['sample_weight', 'scale_pos_weight=50', 'sample_weight

compare_models(3, X_train, y_train, X_test, y_test, model_list, model_names
```

Sample weights are used!

-----

Sample weights are used!

-----


Sample weights are used!

-----

Out[13]:

precision	recall	f1	accuracy	auc
-----------	--------	----	----------	-----

Params	precision	recall	f1	accuracy	auc
Params					
sample_weight Train	0.935	1.000	0.966	0.997	1.000
Test	0.691	0.607	0.647	0.966	0.920
scale_pos_weight=50 Train	0.892	1.000	0.943	0.994	1.000
Test	0.696	0.664	0.679	0.968	0.932
sample_weight + scale_pos_weight=4.5 Train	0.815	1.000	0.898	0.990	1.000
Test	0.607	0.664	0.634	0.961	0.922
sample_weight + scale_pos_weight=50 Train	0.586	1.000	0.739	0.967	1.000
Test	0.411	0.692	0.516	0.934	0.911





```

In [12]: #Compare balance methods at max_depth=3

xgbParams1 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 3,
}

mod1 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams1, 'X

xgbParams2 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 3,
    'scale_pos_weight': 20,
}

mod2 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams2, 'X

xgbParams3 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'max_depth': 3,
    'scale_pos_weight': 1,
}

mod3 = xgb_model_report(3, X_train, y_train, X_test, y_test, xgbParams3, 'X

model_list = [mod1, mod2, mod3]

model_names_list = ['sample_weight', 'scale_pos_weight=50', 'sample_weight

compare_models(3, X_train, y_train, X_test, y_test, model_list, model_names

```

Sample weights are used!

-----


Sample weights are used!

-----

Out[12]:

		precision	recall	f1	accuracy	auc
Params						
sample_weight Train		0.716	1.000	0.834	0.982	0.999
	Test	0.574	0.692	0.627	0.958	0.933
scale_pos_weight=50 Train		0.720	1.000	0.837	0.982	0.999
	Test	0.570	0.682	0.621	0.958	0.928
sample_weight + scale_pos_weight=1 Train		0.716	1.000	0.834	0.982	0.999
	Test	0.574	0.692	0.627	0.958	0.933





### Comments

- The optimum imbalance depends on the max\_depth.
  - max\_depth=6: sample\_weight + scale\_pos\_weight=20
  - max\_depth=5: sample\_weight + scale\_pos\_weight=20
  - max\_depth=4: sample\_weight + scale\_pos\_weight=4.5
  - max\_depth=3: sample\_weight + scale\_pos\_weight=1
- The optimum value may be different when another data set is used.

In [ ]: