

Analysis on All Data Files

On this notebook, I will apply the best optimized models on all five datasets.

Load Libraries

```
In [11]: # Import base libraries
import pandas as pd
import numpy as np
from scipy.io import arff

import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
from xgboost import XGBClassifier

from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import roc_auc_score, roc_curve, auc
from sklearn.metrics import precision_score, recall_score, accuracy_score,
from sklearn.utils import class_weight

from functions import *

from datetime import datetime

import warnings
warnings.filterwarnings('ignore')
```

Load Data

There are five data files:

- * data1, 1year.arff
- * data2, 2year.arff
- * data3, 3year.arff
- * data4, 4year.arff
- * data5, 5year.arff

Note: No cleaning applied to data. XGBoost Classifier can handle the missing values and outliers.

```
In [2]: # Load all five data files

data1 = arff.loadarff('data/1year.arff')
df1 = pd.DataFrame(data1[0])

data2 = arff.loadarff('data/2year.arff')
df2 = pd.DataFrame(data2[0])

data3 = arff.loadarff('data/3year.arff')
df3 = pd.DataFrame(data3[0])

data4 = arff.loadarff('data/4year.arff')
df4 = pd.DataFrame(data4[0])

data5 = arff.loadarff('data/5year.arff')
df5 = pd.DataFrame(data5[0])
```

```
In [3]: # Convert class/label type to binary

df1['class'] = df1['class'].astype('int64')
df2['class'] = df2['class'].astype('int64')
df3['class'] = df3['class'].astype('int64')
df4['class'] = df4['class'].astype('int64')
df5['class'] = df5['class'].astype('int64')
```

```
In [4]: # Size of datasets

print('Size of datasets')
print("Data 1 (Year1):", len(df1))
print("Data 2 (Year2):", len(df2))
print("Data 3 (Year3):", len(df3))
print("Data 4 (Year4):", len(df4))
print("Data 5 (Year5):", len(df5))
```

```
Size of datasets
Data 1 (Year1): 7027
Data 2 (Year2): 10173
Data 3 (Year3): 10503
Data 4 (Year4): 9792
Data 5 (Year5): 5910
```

Imbalance Information

I am using both `sample_weight` and `scale_pos_weight` parameters to deal with the class imbalance.

- `sample_weight`: The weights for training sample are calculated for each dataset separately and used when during training.
- `scale_pos_weight`: I provide certain values to initiate the classifier. I either use the imbalance ratio or square root of the imbalance ratio. These values are not exactly same for the datasets, but close enough to use a constant about average number.
 - Model 7 (max_depth=5): `scale_pos_weight=20` (~imbalance ratio)
 - Model 8 (max_depth=4): `scale_pos_weight=4.5` (~square root of imbalance ratio)

- Model 9 (max_depth=6): scale_pos_weight=20 (~imbalance ratio)

```
In [5]: # Imbalance info using class_weight.compute_class_weight

print('Class Weights:')

df_list = [df1, df2, df3, df4, df5]

for i, df in enumerate(df_list, start=1):
    class_weights = class_weight.compute_class_weight(class_weight='balance
    ratio = class_weights[1]/class_weights[0]
    sqrt_ratio = np.sqrt(class_weights[1]/class_weights[0])
    print(f'Data {i}: Ratio={round(ratio,3)}, sqrt(ratio)={round(sqrt_ratio,3)}')

# The values are very similar for train/test/whole datasets.
# Training weights are used for data training.
```

```
Class Weights:
Data 1: Ratio=24.93, sqrt(ratio)=4.993, class_weights=[ 0.52005625 12.964
94465]
Data 2: Ratio=24.432, sqrt(ratio)=4.943, class_weights=[ 0.52046455 12.71
625   ]
Data 3: Ratio=20.218, sqrt(ratio)=4.496, class_weights=[ 0.52473022 10.60
909091]
Data 4: Ratio=18.014, sqrt(ratio)=4.244, class_weights=[0.52775682 9.5067
9612]
Data 5: Ratio=13.415, sqrt(ratio)=3.663, class_weights=[0.53727273 7.2073
1707]
```

```
In [6]: # Imbalance info using class value_counts

print('Imbalance Ratio, based on class value_counts:')

df_list = [df1, df2, df3, df4, df5]

for i, df in enumerate(df_list, start=1):
    val_counts = df['class'].value_counts()
    ratio= val_counts[0]/val_counts[1]
    sqrt_ratio= np.sqrt(val_counts[0]/val_counts[1])
    print(f'Data {i}: Ratio={round(ratio,3)}, sqrt(ratio)={round(sqrt_ratio,3)}')

# The values are very similar for train/test/whole datasets.
```

```
Imbalance Ratio, based on class value_counts:
Data 1: Ratio=24.93, sqrt(ratio)=4.993
Data 2: Ratio=24.432, sqrt(ratio)=4.943
Data 3: Ratio=20.218, sqrt(ratio)=4.496
Data 4: Ratio=18.014, sqrt(ratio)=4.244
Data 5: Ratio=13.415, sqrt(ratio)=3.663
```

Compare Model 7, 9, 10 Performance on All Datasets

```

In [7]: # Model 9: Best Model for Data 3, max_depth=4

xgbParams_m9 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'scale_pos_weight': 4.5,
    'n_estimators': 125,
    'max_depth': 4,
    'min_child_weight': 3,
    'gamma': 0,
    'learning_rate': 0.20,
    'max_delta_step': 0,
    'reg_lambda': 1,
    'reg_alpha': 0,
    'subsample': 1,
    'colsample_bytree': 1
}

df_list = [df1, df2, df3, df4, df5]
#df_list = [df1, df2]

model_9_df = compare_datafiles_perf(df_list, xgbParams_m9, 'model_9', 1, 1,
model_9_df

```

```


-----model_9-----
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!

```

Out[7]:

	Sample	precision	recall	f1	accuracy	auc
Data						
Data 1	Train	0.9520	1.0000	0.9760	0.9980	1.0000
Data 1	Test	0.7590	0.8040	0.7810	0.9840	0.9630
Data 2	Train	0.8590	1.0000	0.9240	0.9930	1.0000
Data 2	Test	0.4890	0.6340	0.5520	0.9640	0.9240
Data 3	Train	0.7190	1.0000	0.8360	0.9820	1.0000
Data 3	Test	0.5540	0.7200	0.6260	0.9560	0.9160
Data 4	Train	0.7440	1.0000	0.8530	0.9820	1.0000
Data 4	Test	0.4960	0.6600	0.5670	0.9470	0.9420
Data 5	Train	0.9500	1.0000	0.9740	0.9960	1.0000
Data 5	Test	0.6600	0.7330	0.6950	0.9510	0.9580
Average	Train	0.8448	1.0000	0.9126	0.9902	1.0000
Average	Test	0.5916	0.7102	0.6442	0.9604	0.9406

ROC Curve, Dataset Comparison for model_9



```

In [8]: # Model 7, max_depth=5

xgbParams_m7 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'scale_pos_weight': 20,
    'n_estimators': 125,
    'max_depth': 5,
    'min_child_weight': 3,
    'gamma': 0,
    'learning_rate': 0.20,
    'max_delta_step': 0,
    'reg_lambda': 0,
    'reg_alpha': 5,
    'subsample': 1,
    'colsample_bytree': 0.7
}

df_list = [df1, df2, df3, df4, df5]

model_7_df = compare_datafiles_perf(df_list, xgbParams_m7, 'model_7', 1, 1,
model_7_df


-----model_7-----
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!

```

Out[8]:

	Sample	precision	recall	f1	accuracy	auc
Data						
Data 1	Train	0.8800	1.0000	0.9360	0.9950	1.0000
Data 1	Test	0.6270	0.8240	0.7120	0.9760	0.9690
Data 2	Train	0.8730	1.0000	0.9320	0.9940	1.0000
Data 2	Test	0.4840	0.6340	0.5490	0.9640	0.9320
Data 3	Train	0.7700	1.0000	0.8700	0.9860	1.0000
Data 3	Test	0.5380	0.7200	0.6160	0.9540	0.9270
Data 4	Train	0.8060	1.0000	0.8930	0.9870	1.0000
Data 4	Test	0.5560	0.6800	0.6110	0.9550	0.9490
Data 5	Train	0.8290	1.0000	0.9070	0.9860	1.0000
Data 5	Test	0.6070	0.7890	0.6860	0.9450	0.9600
Average	Train	0.8316	1.0000	0.9076	0.9896	1.0000
Average	Test	0.5624	0.7294	0.6348	0.9588	0.9474

ROC Curve, Dataset Comparison for model_7



```
In [9]: # Model 10, max_depth=6
```

```
xgbParams_m10 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'scale_pos_weight': 20,
    'n_estimators': 80,
    'max_depth': 6,
    'min_child_weight': 3,
    'gamma': 0,
    'learning_rate': 0.25,
    'max_delta_step': 4,
    'reg_lambda': 1,
    'reg_alpha': 0,
    'subsample': 1,
    'colsample_bytree': 1,
}

df_list = [df1, df2, df3, df4, df5]

model_10_df = compare_datafiles_perf(df_list, xgbParams_m10, 'model_10', 1,
model_10_df
```


-----model_10-----

Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!

```
Out[9]:
```

	Sample	precision	recall	f1	accuracy	auc
Data						
Data 1	Train	0.9780	1.0000	0.9890	0.9990	1.0000
Data 1	Test	0.6900	0.7840	0.7340	0.9790	0.9560
Data 2	Train	0.9510	1.0000	0.9750	0.9980	1.0000
Data 2	Test	0.5810	0.6060	0.5930	0.9710	0.9060
Data 3	Train	0.9000	1.0000	0.9470	0.9950	1.0000
Data 3	Test	0.6360	0.7010	0.6670	0.9640	0.9310
Data 4	Train	0.8980	1.0000	0.9460	0.9940	1.0000
Data 4	Test	0.5640	0.6410	0.6000	0.9550	0.9480
Data 5	Train	0.9700	1.0000	0.9850	0.9980	1.0000
Data 5	Test	0.6370	0.7220	0.6770	0.9480	0.9510
Average	Train	0.9394	1.0000	0.9684	0.9968	1.0000
Average	Test	0.6216	0.6908	0.6542	0.9634	0.9384

ROC Curve, Dataset Comparison for model_10



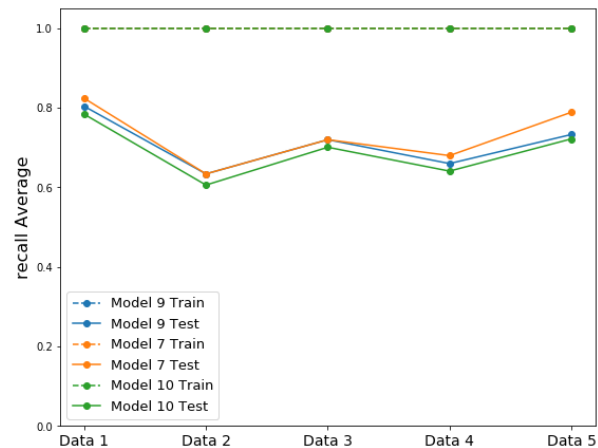



```
In [10]: # Compare models
```

```
model_df_list = [model_9_df, model_7_df, model_10_df]
```

```
model_names_list = ['Model 9', 'Model 7', 'Model 10']
```

```
plot_compare_model_metricsAvg(model_df_list, model_names_list, 1)
```



Comments

Overall the performance of the models are similar, especially recall.

Model 9 (max_depth=4):

- * Previously Selected performance for data 3
- * Overfitting is least on Data 3

Model 7 (max_depth=5):

- * Performance is similar to Model 9
- * Recall is slightly larger than Model 9
- * More smooth performance on all datasets

Model 10 (max_depth=6):

- * Highest metrics
- * But overfitting is larger

Selected model: Model 7

Some Tuning and Comparing

Try some more tuning...

```
In [11]: # sample_weight only, no scale_pos_weight

xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
    #'scale_pos_weight': 4.5,
    'n_estimators': 125,
    'max_depth': 4,
    'min_child_weight': 3,
    'gamma': 0,
    'learning_rate': 0.20,
    'max_delta_step': 0,
    'reg_lambda': 1,
    'reg_alpha': 0,
    'subsample': 1,
    'colsample_bytree': 1
}

df_list = [df1, df2, df3, df4, df5]

model_noSpw_df = compare_datafiles_perf(df_list, xgbParams, 'noSpw', 1, 0,
model_noSpw_df
```

```
-----noSpw-----
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!
```

Out[11]:

	Sample	precision	recall	f1	accuracy	auc
Data						
Data 1	Train	0.9780	1.0000	0.9890	0.9990	1.0000
Data 1	Test	0.8720	0.8040	0.8370	0.9890	0.9660
Data 2	Train	0.9590	1.0000	0.9790	0.9980	1.0000
Data 2	Test	0.6450	0.5630	0.6020	0.9740	0.9400
Data 3	Train	0.8860	1.0000	0.9390	0.9940	1.0000
Data 3	Test	0.6670	0.6360	0.6510	0.9650	0.9290
Data 4	Train	0.9260	1.0000	0.9610	0.9960	1.0000
Data 4	Test	0.6090	0.5150	0.5580	0.9570	0.9450
Data 5	Train	0.9820	1.0000	0.9910	0.9990	1.0000
Data 5	Test	0.7590	0.7330	0.7460	0.9620	0.9630
Average	Train	0.9462	1.0000	0.9718	0.9972	1.0000
Average	Test	0.7104	0.6502	0.6788	0.9694	0.9486

```

In [12]: #n_estimators=100

xgbParams_11 = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'scale_pos_weight': 4.5,
    'n_estimators': 100,
    'max_depth': 4,
    'min_child_weight': 3,
    'gamma': 0,
    'learning_rate': 0.20,
    'max_delta_step': 0,
    'reg_lambda': 1,
    'reg_alpha': 0,
    'subsample': 1,
    'colsample_bytree': 1
}

df_list = [df1, df2, df3, df4, df5]

model_nEst100_df = compare_datafiles_perf(df_list, xgbParams_11, 'nEst100',
model_nEst100_df

```

```

-----nEst100-----
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!

```

Out[12]:

	Sample	precision	recall	f1	accuracy	auc
Data						
Data 1	Train	0.8870	1.0000	0.9400	0.9950	1.000
Data 1	Test	0.6450	0.7840	0.7080	0.9770	0.964
Data 2	Train	0.6970	1.0000	0.8210	0.9820	1.000
Data 2	Test	0.4230	0.6620	0.5160	0.9570	0.923
Data 3	Train	0.6090	1.0000	0.7570	0.9700	1.000
Data 3	Test	0.4910	0.7480	0.5930	0.9480	0.916
Data 4	Train	0.6250	1.0000	0.7690	0.9680	1.000
Data 4	Test	0.4270	0.6800	0.5240	0.9350	0.938
Data 5	Train	0.8420	1.0000	0.9140	0.9870	1.000
Data 5	Test	0.6110	0.7330	0.6670	0.9440	0.959
Average	Train	0.7320	1.0000	0.8402	0.9804	1.000
Average	Test	0.5194	0.7214	0.6016	0.9522	0.940

```

In [13]: # reg_lambda=5

xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'scale_pos_weight': 4.5,
    'n_estimators': 125,
    'max_depth': 4,
    'min_child_weight': 3,
    'gamma': 0,
    'learning_rate': 0.20,
    'max_delta_step': 0,
    'reg_lambda': 5,
    'reg_alpha': 0,
    'subsample': 1,
    'colsample_bytree': 1
}

df_list = [df1, df2, df3, df4, df5]

model_lambda5_df = compare_datafiles_perf(df_list, xgbParams, 'lambda5', 1,
model_lambda5_df

```

```

-----lambda5-----
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!

```

Out[13]:

	Sample	precision	recall	f1	accuracy	auc
Data						
Data 1	Train	0.9090	1.0000	0.9520	0.9960	1.0000
Data 1	Test	0.6890	0.8240	0.7500	0.9800	0.9660
Data 2	Train	0.7850	1.0000	0.8800	0.9890	1.0000
Data 2	Test	0.4800	0.6620	0.5560	0.9630	0.9480
Data 3	Train	0.6470	1.0000	0.7850	0.9750	1.0000
Data 3	Test	0.5070	0.7200	0.5950	0.9500	0.9240
Data 4	Train	0.7400	1.0000	0.8500	0.9810	1.0000
Data 4	Test	0.4770	0.6120	0.5360	0.9440	0.9430
Data 5	Train	0.8600	1.0000	0.9250	0.9890	1.0000
Data 5	Test	0.6190	0.7780	0.6900	0.9470	0.9610
Average	Train	0.7882	1.0000	0.8784	0.9860	1.0000
Average	Test	0.5544	0.7192	0.6254	0.9568	0.9484

```

In [14]: #n_estimators=100
# reg_lambda=5

xgbParams = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'scale_pos_weight': 4.5,
    'n_estimators': 100,
    'max_depth': 4,
    'min_child_weight': 3,
    'gamma': 0,
    'learning_rate': 0.20,
    'max_delta_step': 0,
    'reg_lambda': 5,
    'reg_alpha': 0,
    'subsample': 1,
    'colsample_bytree': 1
}

df_list = [df1, df2, df3, df4, df5]

model_nEst100_lambda5_df = compare_datafiles_perf(df_list, xgbParams, 'nEst
model_nEst100_lambda5_df

```

```

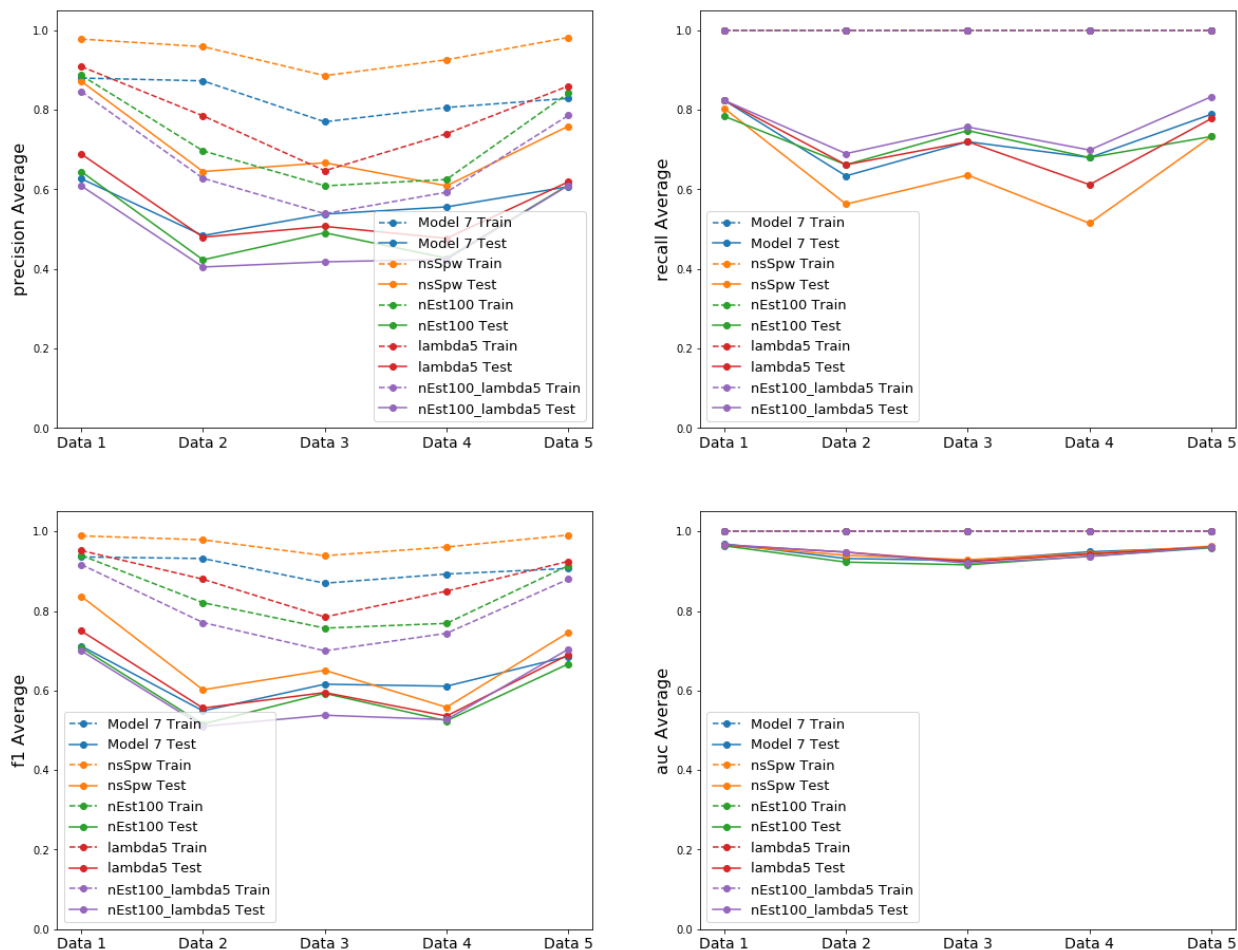
-----nEst100_lambda5-----
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!

```

Out[14]:

	Sample	precision	recall	f1	accuracy	auc
Data						
Data 1	Train	0.8460	1.0000	0.9170	0.9930	1.0000
Data 1	Test	0.6090	0.8240	0.7000	0.9740	0.9670
Data 2	Train	0.6280	1.0000	0.7710	0.9760	1.0000
Data 2	Test	0.4050	0.6900	0.5100	0.9540	0.9480
Data 3	Train	0.5390	1.0000	0.7000	0.9600	1.0000
Data 3	Test	0.4180	0.7570	0.5380	0.9340	0.9200
Data 4	Train	0.5930	1.0000	0.7440	0.9640	1.0000
Data 4	Test	0.4240	0.6990	0.5270	0.9340	0.9370
Data 5	Train	0.7860	1.0000	0.8800	0.9820	1.0000
Data 5	Test	0.6100	0.8330	0.7040	0.9470	0.9610
Average	Train	0.6784	1.0000	0.8024	0.9750	1.0000
Average	Test	0.4932	0.7606	0.5958	0.9486	0.9466

```
In [15]: model_df_list2 = [model_7_df, model_noSpw_df, model_nEst100_df, model_lambda5_df, model_nEst100_lambda5_df]
model_names_list2 = ['Model 7', 'nsSpw', 'nEst100', 'lambda5', 'nEst100_lambda5']
plot_compare_model_metricsAvg(model_df_list2, model_names_list2)
```



Comments

After some tuning, still best performance is Model 7.

Final Model

Model 7 has the best performance.


```
In [16]: # Final Model Parameters
# Model 7, max_depth=5

xgbParams_final = {
    'eval_metric': 'logloss',
    'random_state': 42,
    'scale_pos_weight': 20,
    'n_estimators': 125,
    'max_depth': 5,
    'min_child_weight': 3,
    'gamma': 0,
    'learning_rate': 0.20,
    'max_delta_step': 0,
    'reg_lambda': 0,
    'reg_alpha': 5,
    'subsample': 1,
    'colsample_bytree': 0.7
}
```

```

In [17]: # Final Model Results for All data files

df_list = [df1, df2, df3, df4, df5]


final_model_df = compare_datafiles_perf(df_list, xgbParams_final, 'final_model_df')
final_model_df

-----final_model-----
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!
Sample weights are used!

```

Out[17]:

	Sample	precision	recall	f1	accuracy	auc
Data						
Data 1	Train	0.8800	1.0000	0.9360	0.9950	1.0000
Data 1	Test	0.6270	0.8240	0.7120	0.9760	0.9690
Data 2	Train	0.8730	1.0000	0.9320	0.9940	1.0000
Data 2	Test	0.4840	0.6340	0.5490	0.9640	0.9320
Data 3	Train	0.7700	1.0000	0.8700	0.9860	1.0000
Data 3	Test	0.5380	0.7200	0.6160	0.9540	0.9270
Data 4	Train	0.8060	1.0000	0.8930	0.9870	1.0000
Data 4	Test	0.5560	0.6800	0.6110	0.9550	0.9490
Data 5	Train	0.8290	1.0000	0.9070	0.9860	1.0000
Data 5	Test	0.6070	0.7890	0.6860	0.9450	0.9600
Average	Train	0.8316	1.0000	0.9076	0.9896	1.0000
Average	Test	0.5624	0.7294	0.6348	0.9588	0.9474



```
In [18]: # Data 1 Final Model
d1_final_model = xgb_model_report2(1, df1, xgbParams_final, 'final_model',
```

Sample weights are used!
Data 1 Classification Report:

Training Data:

	precision	recall	f1-score	support
0	1.00	0.99	1.00	5401
1	0.88	1.00	0.94	220
accuracy			0.99	5621
macro avg	0.94	1.00	0.97	5621
weighted avg	1.00	0.99	0.99	5621

Testing Data:

	precision	recall	f1-score	support
0	0.99	0.98	0.99	1355
1	0.63	0.82	0.71	51
accuracy			0.98	1406
macro avg	0.81	0.90	0.85	1406
weighted avg	0.98	0.98	0.98	1406

```
In [19]: # Data 2 Final Model
d2_final_model = xgb_model_report2(2, df2, xgbParams_final, 'final_model',
```

Sample weights are used!

Data 2 Classification Report:

Training Data:

	precision	recall	f1-score	support
0	1.00	0.99	1.00	7809
1	0.87	1.00	0.93	329
accuracy			0.99	8138
macro avg	0.94	1.00	0.96	8138
weighted avg	0.99	0.99	0.99	8138

Testing Data:

	precision	recall	f1-score	support
0	0.99	0.98	0.98	1964
1	0.48	0.63	0.55	71
accuracy			0.96	2035
macro avg	0.74	0.80	0.76	2035
weighted avg	0.97	0.96	0.97	2035

```
In [20]: # Data 3 Final Model
d3_final_model = xgb_model_report2(3, df3, xgbParams_final, 'final_model',
```

Sample weights are used!

Data 3 Classification Report:

Training Data:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	8014
1	0.77	1.00	0.87	388
accuracy			0.99	8402
macro avg	0.88	0.99	0.93	8402
weighted avg	0.99	0.99	0.99	8402

Testing Data:

	precision	recall	f1-score	support
0	0.98	0.97	0.98	1994
1	0.54	0.72	0.62	107
accuracy			0.95	2101
macro avg	0.76	0.84	0.80	2101
weighted avg	0.96	0.95	0.96	2101

```
In [21]: # Data 4 Final Model
d4_final_model = xgb_model_report2(4, df4, xgbParams_final, 'final_model',
```

Sample weights are used!

Data 4 Classification Report:

Training Data:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	7421
1	0.81	1.00	0.89	412
accuracy			0.99	7833
macro avg	0.90	0.99	0.94	7833
weighted avg	0.99	0.99	0.99	7833

Testing Data:

	precision	recall	f1-score	support
0	0.98	0.97	0.98	1856
1	0.56	0.68	0.61	103
accuracy			0.95	1959
macro avg	0.77	0.82	0.79	1959
weighted avg	0.96	0.95	0.96	1959

```
In [22]: # Data 5 Final Model
d5_final_model = xgb_model_report2(5, df5, xgbParams_final, 'final_model',
```

Sample weights are used!

Data 5 Classification Report:

Training Data:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	4408
1	0.83	1.00	0.91	320
accuracy			0.99	4728
macro avg	0.91	0.99	0.95	4728
weighted avg	0.99	0.99	0.99	4728

Testing Data:

	precision	recall	f1-score	support
0	0.98	0.96	0.97	1092
1	0.61	0.79	0.69	90
accuracy			0.95	1182
macro avg	0.79	0.87	0.83	1182
weighted avg	0.95	0.95	0.95	1182

Best Predictor Attributes

```

In [23]: # List important attitudes

final_model_list = [d1_final_model, d2_final_model, d3_final_model, d4_final_model]

df_important_Attr = pd.DataFrame()
cols = df1.columns

for i,df in enumerate(final_model_list, start=1):
    importance = df.feature_importances_
    attribute_importance = pd.DataFrame([cols, importance], index=['Attribute', 'Importance'])
    attribute_importance.sort_values(by='Importance', ascending=False, inplace=True)
    df_important_Attr[f'Data{i}'] = attribute_importance['Attribute'][0:20]

df_important_Attr

```

Out[23]:

	Data1	Data2	Data3	Data4	Data5
0	Attr34	Attr59	Attr52	Attr62	Attr62
1	Attr6	Attr34	Attr34	Attr33	Attr34
2	Attr22	Attr10	Attr10	Attr31	Attr64
3	Attr24	Attr5	Attr26	Attr6	Attr7
4	Attr25	Attr27	Attr27	Attr34	Attr27
5	Attr27	Attr7	Attr59	Attr27	Attr6
6	Attr26	Attr26	Attr31	Attr26	Attr42
7	Attr8	Attr25	Attr16	Attr42	Attr16
8	Attr45	Attr24	Attr35	Attr52	Attr49
9	Attr50	Attr16	Attr24	Attr13	Attr5
10	Attr1	Attr22	Attr25	Attr57	Attr48
11	Attr51	Attr58	Attr20	Attr21	Attr25
12	Attr5	Attr15	Attr13	Attr64	Attr51
13	Attr59	Attr11	Attr41	Attr41	Attr43
14	Attr21	Attr42	Attr5	Attr1	Attr31
15	Attr49	Attr41	Attr37	Attr5	Attr23
16	Attr13	Attr28	Attr45	Attr37	Attr2
17	Attr37	Attr20	Attr58	Attr24	Attr50
18	Attr38	Attr31	Attr19	Attr59	Attr21
19	Attr35	Attr37	Attr6	Attr15	Attr14

```
In [24]: #Common Best predictors
```

```
set1 = set(df_important_Attr['Data1'])
set2 = set(df_important_Attr['Data2'])
set3 = set(df_important_Attr['Data3'])
set4 = set(df_important_Attr['Data4'])
set5 = set(df_important_Attr['Data5'])

important_attr_common = set.intersection(set1, set2, set3, set4, set5)
important_attr_common
```


```
Out[24]: {'Attr27', 'Attr34', 'Attr5'}
```

```
In [25]: # Scatter Graphs, Data 3
```

```
plt.figure(figsize=(20, 5))

for i, col in enumerate(important_attr_common, start=1):
    plt.subplot(1, 3, i)
    plt.scatter(df3[col], df3['class'], alpha=0.1)
    plt.title(col)

plt.savefig('figures/scatter_d3_importantAttr.png')
```




```
In [26]: # Histograms
```

```
plt.figure(figsize=(20, 5))

for i, col in enumerate(important_attr_common, start=1):
    plt.subplot(1, 3, i)
    plt.hist(df3[col], bins=100, log=True)
    plt.title(col)

plt.savefig('figures/hist_d3_log_importantAttr.png')
```




```
In [27]: # Box plots
```

```
plt.figure(figsize=(20, 5))

for i, col in enumerate(important_attr_common, start=1):
    plt.subplot(1, 3, i)
    plt.boxplot(df3[col].dropna())
    plt.title(col)

plt.savefig('figures/boxPlot_d3_importantAttr.png')

## Boxplot only draws the attribute only when missing entries removed.
```



Best predictor descriptions

- X27: profit on operating activities / financial expenses
- X34: operating expenses / total liabilities
- X5: $[(\text{cash} + \text{short-term securities} + \text{receivables} - \text{short-term liabilities}) / (\text{operating expenses} - \text{depreciation})] * 365$

Interpretation of Results

5-year Period (Data 1):

- Model correctly identifies the 80.4% of the true bankrupt companies, which will bankrupt 5 years later. (recall)
- Among the model predicted bankruptcy companies, 64.1% of them are true bankrupt companies, which will bankrupt 5 years later. (precision)
- The Harmonic Mean of Precision and Recall (f1-score) is 71.3%.

4-year Period (Data 2):

- Model correctly identifies the 62.0% of the true bankrupt companies, which will bankrupt 4 years later. (recall)
- Among the model predicted bankruptcy companies, 50.6% of them are true bankrupt companies, which will bankrupt 4 years later. (precision)
- The Harmonic Mean of Precision and Recall (f1-score) is 55.7%.

3-year Period (Data 3):

- Model correctly identifies the 72.0% of the true bankrupt companies, which will bankrupt 3 years later. (recall)
- Among the model predicted bankruptcy companies, 53.5% of them are true bankrupt companies, which will bankrupt 3 years later. * The Harmonic Mean of Precision and Recall (f1-score) is 61.4%.

2-year Period (Data 4):

- Model correctly identifies the 68.0% of the true bankrupt companies, which will bankrupt 2 years later. (recall)
- Among the model predicted bankruptcy companies, 55.6% of them are true bankrupt companies, which will bankrupt 2 years later. (precision)
- The Harmonic Mean of Precision and Recall (f1-score) is 61.1%.

1-year Period (Data 5):

- Model correctly identifies the 78.9% of the true bankrupt companies, which will bankrupt 1 years later. (recall)
- Among the model predicted bankruptcy companies, 60.7% of them are true bankrupt companies, which will bankrupt 1 years later. (precision)
- The Harmonic Mean of Precision and Recall (f1-score) is 68.6%.

On Average:

- Model correctly identifies the 72.3% of the true bankrupt companies. (recall)
- Among the model predicted bankruptcy companies, 56.9% of them are true bankrupt companies. (precision)
- The Harmonic Mean of Precision and Recall (f1-score) is 63.6%.

Class 0 predictions:

- Model correctly identifies the ~97% of the true still operating companies. (recall, class 0)

- Among the model predicted still operating companies, ~98% of them are true still operating companies. (precision, class 0)
- The Harmonic Mean of Precision and Recall (f1-score, class 0) is ~98%.

Conclusion

I had three main challenges in this project:

1. Class Imbalance:

- Resolved.
- Balancing the sample increased the recall and so decreased the precision. The balance sample has moderate evaluation scores. See Model 2 for Data 3.
- However overfitting didn't improve much.

2. Low recall score:

- I managed to improve recall significantly for all datasets. For instance, Data 3 recall score improved from 0.467 to 0.720.
- I couldn't go for higher recall value, because the precision was decreasing dramatically (below 0.5). The recall and precision are inversely proportional.

3. Large overfitting:

- I did decrease the overfitting, but not on desired level.
- I tuned parameters which are effective on overfitting, and find the optimum designs that produces low overfitting, large recall and moderate precision.
- However, I couldn't enforce larger reduction in overfitting, since it causes the precision go below 0.5. which is the random guess probability.

Overall, my model correctly identifies

- 72.3% of the true bankrupt companies
- 97% of the true still operating companies

Future Work

- Create separate final models for each dataset; not just one final model that applies on all.
- Search for alternative Classifier methods/tools.
- Simply/shorten functions that are created during the project. They have repeating codes.